Adrian Ng MSc.

Seeking Junior-Level Data Engineering Opportunities

Email: contact@adrian.ng
Website: adrian.ng

Profile

I am a Computer Science graduate passionate about Data Engineering and I seek opportunities that meet my growing experience in *Java* – a language I have used in numerous academic projects ranging from the implementation of financial models to large-scale data processing with *Apache Hadoop MapReduce*.
Prior to postgraduate study, my expertise was in *SQL development* focusing on the implementation of segmentation processes for a number of clients including: *Virgin Media*, *TUI*, *UPC*, *MSD*, *Volkswagen*, *KwikFit*.
After graduation, my most recent accomplishments as a Data Analyst at *Manchester City FC* were in the technical parts (e.g. pipelines, architecture) of the projects I worked on, which leads me now to pursue a career in Data Engineering.

Education

- **Master of Science in Data Science and Analytics** — with Distinction
  Department of Computer Science, Royal Holloway — *Sept. 2016 – Dec. 2017*
  Java Modules: ○ Programming for Data Analysis ○ Large Scale Data Storage and Processing ○ Methods of Computational Finance ○ Dissertation
- **Bachelor of Engineering in Mechanical Engineering** — Upper Second Class with Honours
  School of Engineering, King's College London — *Sept. 2007 – July 2010*

Technologies

Languages • Java 8 • SQL
Software • IntelliJ IDEA • SQL Server Management Studio • Git • Jira • Maven

Java Projects

- **Implementation of Value at Risk (VaR) measures in Java** (https://adrian.ng/java/var/) (https://github.com/Adrian-Ng/VaR)
  Assuming a number of hypothetical investment portfolios, my dissertation project implemented a number of approaches to estimating *VaR*, a measure of risk, and variance/volatility (for model parameterization).
  - **VaR:** a) *Model Building* b) *Historical Simulation* c) *Monte Carlo Simulation*.
  - **Variance/Volatility:** a) *Equal Weighted* b) *Exponentially Weighted Moving Average (EWMA)* c) *GARCH(1,1)*.
    ∗ *EWMA* lambda parameter taken from J.P. Morgan's *RiskMetrics*. ∗ *GARCH(1,1)* parameters were found via *Levenberg-Marquardt* optimisation.

  Because of these numerous approaches, object-oriented techniques and patterns were implemented. In addition, I used Java's concurrency APIs to parallelize the 100,000+ random walks generated by *Monte Carlo* for simulating stock prices. Real world financial data was obtained via *Google Finance* and *Yahoo Finance APIs*.
- **Option Pricing** (https://adrian.ng/java/options/) (https://github.com/Adrian-Ng/OptionPricer)
  This project implements three approaches to estimating option prices in Java:
  ○ Monte Carlo Simulation ○ Black Scholes ○ Binomial Trees
  Apache Commons Math API was used to deal with some probabilistic assumptions.
- **Data Mining with Hadoop MapReduce** (https://github.com/Adrian-Ng/HadoopEnron)
  I wrote number of *MapReduce* applications in Java including extracting the communications network from the *Enron Corpus*, a large dataset of emails, or aggregation of Twitter data.
  Applications were exported and executed on *Hadoop* clusters (both single node and distributed). Input/Output datasets were stored in HDFS and accessed via `hadoop fs` commands.
  A subsequent exercise was undertaken to minimise the verbosity of these *Hadoop MapReduce* applications by translating them to *Scala* for use in a *Spark REPL*.
- **Java 8 Streams with financial data** (https://adrian.ng/java/yahoofinance/#stream)
  A small exercise involving the use of *Java 8 Streams*. Processing real-world financial data to return *mean* and *variance* of some market asset.

Manchester City Football Club

Data Analyst *Fan Relationship Management* *Jan. - July 2018*

- **New York City FC Project:** I took ownership of this project to integrate *NYCFC's* transactional and demographic data with *City Football Group's* data-warehouse. This six-month project involved many phases including: discovery, engineering, and analysis. Data came from multiple external sources each with differing schema: *NYCFC*, *Ticketmaster Salesforce*, *Major League Soccer*.
  - **Data Pipeline:** I implemented a data pipeline to ingress data from a number of remote *SQL* databases. This process was encapsulated in *stored procedures* which used appropriate DML & DDL (`OPENQUERY`, `MERGE`) for efficient ETL. This pipeline replaced the slower front-end *Informatica* solution.
  - **Data Cubes:** I used an aggregated dataset to compare the distribution of `NULL` values. These analyses were transformed to *Data Cubes* to pre-compute every possible roll-up/drill-down. As such, bandwidth was minimised across our distributed servers and need for real-time computation in *Tableau* front-end was eliminated, resulting in improved user-experience.
  - **Mentoring:** As part of this project, I was dedicated to mentoring a junior colleague remotely in New York. I organised weekly workshops to teach basic DML and more advanced DDL with a goal toward self-sufficiency in writing database queries and working with stored procedures. Additional material on my website helped supplement these workshops.
- **GDPR Stream Integration:** I worked on the integration of a GDPR preference stream into our data stores (*SQL*, *Salesforce*). I implemented a new pipeline and refactored numerous processes downstream . I worked with the development team to provided specification and UAT testing. I built an efficient, automated `MERGE` process using primary key constraints, clustered indexes, triggers.
- **Customer Churn Model:** I contributed datasets and collaborated on feature/model selection. In particular, looking at *logistic regression* and *Beta-Geometric/Beta-Bernoulli* models in R Studio.

ITG Creator

Senior CRM Campaign Executive *SQL Development* *Dec. 2013 - Sept. 2016*

The majority of my work in this role involved working with SQL processes which were used to transform customer data into CRM segmentations. As senior team member, I developed a number of these processes. On occasion, I held responsibility for resourcing and managing the team's workload using *Jira*.

- **Virgin Media Segmentation** (https://adrian.ng/SQL/cte/Recursion/ (https://adrian.ng/SQL/misc/openquery-xml)
  I built an end-to-end segmentation process in *SQL*. This included building a fast, flexible, and bespoke import tool around `BULK INSERT`. Remote server queries (`OPENQUERY`) made use of `XML` to effectively `INNER JOIN` local and remote tables resulting in speed and minimial resource use on a busy live server. Recursive queries were used to implement a solution (similar to `flatMap` in *Java 8*) for efficient *regex*.
- **Volkswagen Onboarding:** I worked with `.NET` developers and project managers to bring Volkswagen on-board as a new client. This required implementing a new segmentation process for broadcasting email *and* SMS. In addition, I provided specification to developers for their data warehousing/archiving ingress schema.
- **TUI Redesign:** I collaborated closely with the TUI client during a three-month project to redesign the existing *Thomson* and *First Choice* mailings. `TCL` scripts were developed to merge dynamic content into the `HTML` body. My efforts on this project were awarded by the client.

Seatwave (now Ticketmaster)

Marketing Analyst Intern *Commercial Team* *May 2013 - Dec. 2013*

Using *SQL Server Management Studio* for the first time, I wrote *DML* capable of querying the transactional/customer databases to return data for warehousing, reporting, and segmentation. I also worked on pricing and spatial analyses (QGIS).